
**Graph Neural Networks for Efficient Brain Tumor Segmentation:
A Superpixel-Based Approach with $6.9\times$ Speedup**

By

Sakib Khan	ID: 21225103290
Rifa Sanjida	ID: 21225103274
Kishor Kumar Das	ID: 21225103295
Md. Mahmudul Hasan	ID: 21225103275
Md. Minhajur Rahman	ID: 21224103054

Submitted in partial fulfillment of the requirements of the degree of

**Bachelor of Science in
Computer Science and Engineering**

[scale=0.2]image/BUBT-Logo.png

Department of Computer Science and Engineering
Bangladesh University of Business and Technology

January 2026

Declaration

We do hereby declare that the research works presented in this thesis entitled, “Graph Neural Networks for Efficient Brain Tumor Segmentation: A Superpixel-Based Approach with $6.9\times$ Speedup” are the results of our works. We further declare that the thesis/report has been compiled and written by us, and no part of this thesis/report has been submitted elsewhere for the requirements of any degree, award, diploma, or any other purposes except for publications. The materials that are obtained from other sources are duly acknowledged in this thesis.

Signature of Authors

Sakib Khan

ID: 21225103290

Rifa Sanjida

ID: 21225103274

Kishor Kumar Das

ID: 21225103295

Md. Mahmudul Hasan

ID: 21225103275

Md. Minhajur Rahman

ID: 21224103054

Approval

I do hereby declare that the research works presented in this thesis entitled, “Graph Neural Networks for Efficient Brain Tumor Segmentation: A Superpixel-Based Approach with $6.9\times$ Speedup”, are the outcome of the original works carried out by the Sakib Khan, Rifa Sanjida, Kishor Kumar Das, Md. Mahmudul Hasan, Md. Minhajur Rahman under my supervision. I further declare that no part of this thesis/report has been submitted elsewhere for the requirements of any degree, award, diploma, or other purposes except for publications. I further certify that the dissertation meets the requirements and standards for the degree of Bachelor of Science in Computer Science and Engineering.

Supervisor

Mr. Shamim Ahmed

Assistant Professor

Department of Computer Science & Engineering

Bangladesh University of Business and Technology

Dhaka, Bangladesh

Chairman

Md. Saifur Rahman

Assistant Professor & Chairman

Department of Computer Science & Engineering

Bangladesh University of Business and Technology

Dhaka, Bangladesh

Dedication

We would like to dedicate this research to our loving parents . . .

Acknowledgement

We would like to express our heartfelt gratitude to the almighty Allah who offered our family and us kind care throughout this journey. We are deeply thankful to the Bangladesh University of Business and Technology (BUBT) for providing us with such a wonderful environment to pursue our research. We would like to express our sincere gratitude to Mr. Shamim Ahmed, Assistant Professor, Department of Computer Science and Engineering, BUBT. We have completed our research with his help. We found the research area, topic, and problem with his suggestions. He guided us with our study and supplied us with many research papers and academic resources in this area. He is patient and responsible. When we had questions and needed his help, he would always find time to meet and discuss with us no matter how busy he was. We also want to give thanks to Md. Saifur Rahman, Assistant Professor & Chairman, Department of Computer Science and Engineering, and all our teachers for providing a solid background for our studies and research thereafter. They have been a great source of inspiration to us, and we thank them from the bottom of our hearts. We would also like to acknowledge our team members for supporting each other and be grateful to our university for providing this opportunity for us.

Abstract

Getting accurate brain tumor segmentation is essential for diagnosis and treatment planning, but there’s a problem: current deep learning methods need serious computational resources that make clinical deployment challenging. This thesis explores a different approach—using graph neural networks for efficient binary brain tumor segmentation on the BraTS 2021 dataset. Instead of processing entire 3D MRI volumes, we convert them into graph representations using superpixels, with each node containing 15 engineered features. This lets us keep competitive accuracy while dramatically cutting computational costs.

We validated the approach through 5-fold cross-validation with patient-level stratified splits. Our 5-layer GraphSAGE model hits $90.39\% \pm 0.69\%$ Dice coefficient, and when we ensemble the models, performance jumps to 92.92%. More importantly for practical use, we’re seeing $6.9\times$ faster inference (12.7ms vs. 87.8ms per patient) and $156\times$ fewer parameters (439K vs. 68M) compared to standard 3D U-Net. GPU memory usage? Just 2.1 GB instead of 8.4 GB.

Through systematic testing, we confirmed that 5 layers work best—adding a 6th layer doesn’t help. We also found something interesting: batch size really matters in medical imaging GNN tasks. We ran 15 integrity checks throughout to ensure zero data leakage and full reproducibility with deterministic training.

What this shows is that graph neural networks can deliver a strong accuracy-efficiency trade-off for medical image segmentation. The 92.92% ensemble Dice score combined with the speed and memory improvements makes this practical for real-time clinical use, especially in resource-constrained environments where deploying massive volumetric models just isn’t feasible.

Keywords: Brain Tumor Segmentation, Graph Neural Networks, Medical Image Analysis, BraTS 2021, Computational Efficiency, Deep Learning

List of Abbreviations and Acronyms

3D	Three-Dimensional
ABET	Accreditation Board for Engineering and Technology
AI	Artificial Intelligence
BAETE	Bangladesh Accreditation Council for Engineering and Technical Education
BraTS	Brain Tumor Segmentation
CE	Contrast-Enhancing
CNN	Convolutional Neural Network
CPU	Central Processing Unit
CT	Computed Tomography
CUDA	Compute Unified Device Architecture
CV	Cross-Validation
DICOM	Digital Imaging and Communications in Medicine
FDA	Food and Drug Administration
FLAIR	Fluid-Attenuated Inversion Recovery
GAT	Graph Attention Network
GDPR	General Data Protection Regulation
GE	General Electric
GNN	Graph Neural Network
GPU	Graphics Processing Unit
GraphSAGE	Graph Sample and Aggregate
HIPAA	Health Insurance Portability and Accountability Act

HPC	High-Performance Computing
LGG	Low-Grade Glioma
ML	Machine Learning
MRI	Magnetic Resonance Imaging
NIfTI	Neuroimaging Informatics Technology Initiative
PACS	Picture Archiving and Communication System
ReLU	Rectified Linear Unit
RTX	Ray Tracing Texel eXtreme
SLIC	Simple Linear Iterative Clustering
T1	T1-weighted
T1ce	T1-weighted Contrast-Enhanced
T2	T2-weighted
TCGA	The Cancer Genome Atlas
U-Net	U-shaped Network
VRAM	Video Random Access Memory
WHO	World Health Organization

Contents

Declaration	ii
Approval	iii
Dedication	iv
Acknowledgement	v
Abstract	vi
List of Abbreviations and Acronyms	viii
1 Introduction	1
1.1 Problem Statement	1
1.2 Problem Background	2
1.3 Research Objectives	3
1.4 Motivations	4
1.5 Significance of the Research	5
1.6 Research Contribution	5
1.7 Complex Engineering Analysis	7
1.7.1 Complex Engineering Problems (P1-P7)	7
1.7.2 Complex Engineering Activities (A1-A5)	9
1.8 Thesis/Project Organization	10
2 Background Study and Literature Review	13
2.1 Evolution of Medical Image Segmentation	13
2.1.1 Classical Approaches	13
2.1.2 The CNN Revolution: U-Net and Its Variants	13
2.1.3 Transformer-Based Architectures	14
2.2 Graph Neural Networks in Medical Imaging	14
2.2.1 Fundamentals of Graph Representation Learning	14

2.2.2	Supixel Segmentation	15
2.2.3	Graph-Based Medical Segmentation: Existing Work	15
2.3	Research Gap and Motivation	16
2.3.1	The Efficiency-Accuracy Dilemma	16
2.3.2	Unexplored GNN Potential	16
2.4	Positioning This Work	17
3	Research Methodology	18
3.1	Proposed Framework	18
3.2	Data Set Analysis	20
3.2.1	Data Collection & Preprocessing	21
3.2.2	Feature Engineering and Selection	22
3.3	Algorithm/ Model Analysis	23
3.3.1	Graph Construction Algorithm	23
3.3.2	GraphSAGE Architecture	24
3.3.3	Training Protocol	25
3.3.4	Ensemble Strategy	26
3.3.5	Evaluation Metrics	26
3.4	Implementation Details	27
4	Experimental Results and Analysis	29
4.1	Environment Setup	29
4.2	Cross-Validation Performance	30
4.2.1	Individual Fold Results	31
4.2.2	Secondary Metrics	32
4.3	Ensemble Performance	33
4.3.1	Ensemble Strategy and Results	33
4.4	Efficiency Benchmarking	34
4.4.1	Baseline Comparison	34
4.4.2	Clinical Deployment Implications	35
4.5	Ablation Studies	36

4.5.1	Network Depth Analysis	36
4.5.2	Batch Size Sensitivity	37
4.6	Qualitative Analysis	38
4.7	Key Findings and Discussions	39
4.7.1	Primary Contributions Validated	39
4.7.2	Limitations and Constraints	40
4.8	Comparison with State-of-the-Art Methods	41
5	Standards, Constraints, and Milestones	43
5.1	Sustainability Standards	44
5.1.1	Software Development Standards	44
5.1.2	Data Management Standards	44
5.1.3	Environmental Sustainability	45
5.2	Impacts on Society	45
5.2.1	Healthcare Accessibility	45
5.2.2	Clinical Workflow Integration	46
5.2.3	Diagnostic Equity	46
5.3	Ethics	47
5.3.1	Algorithmic Fairness and Bias	47
5.3.2	Privacy and Data Protection	47
5.3.3	Clinical Accountability	48
5.3.4	Informed Consent	48
5.4	Challenges	49
5.4.1	Technical Challenges	49
5.4.2	Operational Challenges	49
5.4.3	Validation Challenges	50
5.5	Constraints	50
5.5.1	Design Constraints	50
5.5.2	Component Constraints	51
5.5.3	Budget Constraints	51

5.6	Timeline and Gantt Chart	52
6	Conclusion	56
6.1	Limitations	56
6.2	Future Works and Direction	57

Chapter 1

1 Introduction

Segmenting brain tumors from MRI scans is crucial for diagnosis, treatment planning, and tracking how well therapies are working. The BraTS (Brain Tumor Segmentation) challenge has become the go-to benchmark for testing automated segmentation algorithms on multi-parametric MRI data. Binary segmentation—simply detecting tumor vs. non-tumor tissue—might seem straightforward compared to multi-class approaches, but it’s actually essential for screening workflows, quick assessments, and initial diagnostics.

Most current methods rely heavily on volumetric CNNs, especially 3D U-Net architectures. These work well but come with a catch: they’re computationally expensive. We’re talking 60-80 million parameters and needing 8+ GB of GPU memory, which creates real barriers when you want to deploy these in clinics with limited resources, on edge devices, or anywhere you need fast results. Healthcare systems need efficient automated tools, which means we need approaches that stay accurate while cutting down the computational demands significantly.

This is where Graph Neural Networks come in. Instead of processing entire 3D volumes, GNNs represent images as graphs—superpixels become nodes, and spatial relationships become edges. This lets them capture local texture and global structure while working with much smaller data. It’s actually a natural fit for medical images, and it takes advantage of an important fact: tumors typically only occupy 5-10% of brain volume in MRI scans, so there’s inherent sparsity we can exploit.

1.1 Problem Statement

Here’s the issue: automated brain tumor segmentation is clinically important, but current deep learning methods are stuck in an efficiency-accuracy trade-off. Volumetric CNNs hit high accuracy numbers, sure, but they need massive computational

resources—68M+ parameters, over 8 GB GPU memory, and 80+ ms per patient inference time. That makes them impractical for most real-world clinical settings. On the flip side, lightweight alternatives usually sacrifice too much accuracy to be useful.

So the core question becomes: *Can we build a segmentation method that hits competitive accuracy ($>90\%$ Dice) while still being efficient enough for resource-constrained clinics?*

Answering this means tackling several technical hurdles. First, we need to convert 3D medical images into meaningful graph representations without losing the diagnostic features that matter. Second, the neural architecture has to be compact but still expressive enough to detect tumors reliably. Third, validation needs to be rigorous—absolutely zero data leakage. And finally, we need to show real efficiency gains through solid benchmarking.

1.2 Problem Background

The BraTS (Brain Tumor Segmentation) challenge, established in 2012, has become the de facto benchmark for evaluating brain tumor segmentation algorithms. The BraTS 2021 dataset comprises 1,251 multi-institutional glioma cases, each with four MRI modalities (T1, T1 contrast-enhanced, T2, FLAIR) and expert-annotated ground truth segmentations. The challenge focuses on three nested tumor sub-regions: enhancing tumor (ET), tumor core (TC), and whole tumor (WT).

Traditional approaches relied on handcrafted features and classical machine learning methods (SVM, Random Forests), achieving Dice scores around 70-80%. The deep learning revolution, particularly 3D U-Net architectures introduced in 2015, improved performance to 85-90% Dice for state-of-the-art systems. Recent winners of BraTS challenges typically employ ensemble methods combining multiple volumetric CNNs, achieving 88-92% Dice but at the cost of even greater computational requirements.

Graph Neural Networks emerged as a promising alternative in the broader computer vision community, with successful applications in social network analysis,

molecular property prediction, and recommendation systems. However, their application to medical imaging remains limited, with only a handful of preliminary studies showing feasibility but often lacking rigorous validation or efficiency analysis. Critical questions remain about optimal graph construction strategies, feature engineering for medical data, and whether GNNs can match CNN performance while delivering on their efficiency promises.

1.3 Research Objectives

The primary objective of this research is to develop and rigorously validate an efficient graph-based neural network approach for brain tumor segmentation that achieves competitive accuracy while providing substantial computational advantages. Specific objectives include:

- Develop a robust graph construction pipeline that converts 3D MRI volumes into 2D graph representations, preserving critical tumor features while reducing computational complexity through superpixel-based node generation
- Design and optimize a graph neural network architecture specifically for binary brain tumor segmentation, balancing model capacity with parameter efficiency
- Conduct rigorous 5-fold cross-validation with patient-level stratified splits to ensure generalization capability and statistical validity, implementing comprehensive data integrity checks
- Perform extensive efficiency analysis comparing inference speed, memory consumption, and model size against established U-Net baseline to quantify practical deployment advantages
- Execute systematic ablation studies to validate architectural design choices (network depth, hidden dimensions, batch size) and identify optimal configurations
- Implement ensemble methods to explore upper performance bounds and demonstrate robustness through model averaging strategies

- Ensure complete reproducibility through deterministic training protocols, seed management, and comprehensive auditing procedures (targeting 15/15 integrity checks)
- Generate qualitative visualizations and quantitative analyses suitable for clinical interpretation and publication-quality reporting

1.4 Motivations

A few things pushed this research in this direction:

Clinical Deployment Challenges: Talking with radiologists and clinical partners, it became clear that even though current deep learning models are accurate, they’re hard to actually deploy. The hardware requirements, slow inference, and complexity of integrating with existing hospital systems create real problems.

Resource-Constrained Settings: Lots of healthcare facilities—especially in developing regions—don’t have access to the high-end GPUs that volumetric CNNs need. This creates an equity problem where AI-assisted diagnosis is only available to well-funded institutions.

Real-Time Applications: New clinical applications like intra-operative guidance, emergency screening, and telemedicine need near-instant results. Current methods just can’t deliver that consistently.

Theoretical Promise of Graphs: Tumors are sparse—they typically only take up 5-10% of brain volume. GNNs have worked well in other domains, so it seemed like graph representations might be able to exploit this sparsity for computational efficiency.

Data Integrity Concerns: There’ve been concerning reports of data leakage in medical imaging studies lately. This motivated us to implement comprehensive auditing procedures (we ended up with 15/15 integrity checks passing).

Efficiency-Accuracy Trade-off: The broader AI community has been focusing more on sustainable and efficient AI. We wanted to see if medical imaging could benefit from similar improvements without losing clinical usefulness.

1.5 Significance of the Research

This research makes significant contributions to both the medical imaging and machine learning communities:

Clinical Impact: By achieving 92.92% Dice with $6.9\times$ faster inference and $156\times$ parameter reduction, our approach enables deployment in resource-constrained clinical settings, emergency departments, and mobile diagnostic units where current methods are impractical.

Methodological Advancement: We demonstrate that graph-based approaches can match volumetric CNN performance on medical segmentation tasks, challenging the prevailing assumption that 3D convolutions are necessary for medical image analysis.

Efficiency Paradigm: Our work contributes to the growing body of evidence that efficiency and accuracy need not be mutually exclusive, particularly relevant as healthcare AI scales to billions of global users.

Validation Rigor: The comprehensive integrity validation framework (15/15 checks including feature dimension verification, patient leakage detection, label validity, and normalization checks) sets a higher standard for medical AI research reproducibility.

Ablation Insights: Discovery of batch size sensitivity in medical imaging (batch 32 optimal, degrading to 83% at batch 64) and architectural validation (5 layers = 6 layers, no benefit from added depth) provides actionable guidance for future medical imaging GNN designs.

Open Science: Complete pipeline documentation, reproducible code with deterministic training (seed 42), and detailed methodology enable other researchers to build upon this work and validate our findings.

1.6 Research Contribution

This thesis makes the following specific contributions:

- **Novel Graph Construction Pipeline:** A superpixel-based graph genera-

tion method using SLIC algorithm with 200 adaptive slices per patient and 15 engineered features (intensity statistics, spatial information, shape/texture descriptors) that captures tumor characteristics without ground-truth leakage.

- **Optimized GNN Architecture:** A 5-layer GraphSAGE model with 256 hidden dimensions and 439,041 parameters, validated through systematic ablation showing no benefit from deeper networks (6 layers: 84.00% vs. 5 layers: 84.03%).
- **Strong Performance with Efficiency:** $90.39\% \pm 0.69\%$ Dice on 5-fold cross-validation and 92.92% with ensemble, while achieving $6.9\times$ faster inference (12.7ms vs. 87.8ms) and $156\times$ parameter reduction (439K vs. 68M) compared to 3D U-Net.
- **Comprehensive Validation Framework:** Patient-level stratified 5-fold splits with zero data leakage, verified through paranoid auditing (15/15 integrity checks: feature dimensions, patient overlap, label validity, normalization, seed consistency, batch configuration).
- **Batch Size Discovery:** Empirical finding that smaller batches (32) significantly outperform larger batches (64: 83%, 48: 86%) in medical imaging GNN tasks, providing actionable insights for future research.
- **Complete Reproducibility Package:** Deterministic training protocols (seed 42), comprehensive pipeline documentation (8-phase guide), and validated code enabling full experimental reproduction.
- **Clinical Deployment Roadmap:** Demonstration that real-time inference (12.7ms), low memory footprint (2.1 GB), and small model size (439K parameters) enable deployment on edge devices, mobile platforms, and resource-constrained clinical environments.

1.7 Complex Engineering Analysis

This research addresses complex engineering challenges requiring depth of knowledge across graph neural networks, medical image processing, and clinical validation. The work navigates conflicting requirements between accuracy (target >90% Dice) and efficiency (target <50ms inference, <1M parameters), conducts extensive analysis through stratified 5-fold cross-validation with 15 integrity checks, and tackles domain-specific issues including MRI protocol variability, extreme class imbalance (tumors occupy 5-10% of brain volume), and zero data leakage requirements critical for valid medical AI.

The project integrates diverse resources (BraTS 2021 dataset, PyTorch Geometric, SimpleITK, consumer-grade RTX 2060 6GB GPU) and requires interdisciplinary interaction bridging computer science, medical imaging, and clinical practice. Novel contributions include the first comprehensive GNN efficiency study for brain segmentation, discovery of batch size sensitivity in medical imaging GNN tasks, and demonstration that graph-based methods achieve CNN-competitive accuracy (92.92%) with $6.9\times$ speedup and $156\times$ parameter reduction. These efficiency gains enable deployment in resource-constrained settings (developing countries, rural clinics), directly impacting healthcare accessibility and contributing to sustainable AI through reduced energy consumption at scale.

1.7.1 Complex Engineering Problems (P1-P7)

This work addresses seven categories of complex engineering problems as defined by accreditation standards (ABET, BAETE):

P1: Depth of Knowledge Required - Advanced understanding across graph neural networks (GraphSAGE, GAT, message-passing), medical image processing (SLIC superpixels, 3D MRI preprocessing, multi-modal intensity analysis), deep learning optimization (Adam, batch normalization, learning rate scheduling), and software engineering (PyTorch Geometric, CUDA optimization, NIfTI/DICOM standards, version control).

P2: Range of Conflicting Requirements - Navigation of competing objectives: accuracy vs. speed ($>90\%$ Dice while maintaining $<50\text{ms}$ inference), expressiveness vs. efficiency (tumor detection within $<1\text{M}$ parameters), feature richness vs. construction speed (15 engineered features without excessive overhead), and validation rigor vs. iteration speed (5-fold cross-validation with integrity checks).

P3: Depth of Analysis Required - Extensive experimental validation through stratified 5-fold cross-validation with patient-level splits (900 train, 100 validation, 251 test per fold), systematic ablation studies (5 vs. 6 layers, batch sizes 32/48/64), hyperparameter tuning across 20+ configurations, comprehensive integrity auditing (15 validation checks ensuring zero data leakage), and multi-metric evaluation (Dice, sensitivity, specificity, precision, inference time, memory, parameters).

P4: Familiarity with Domain-Specific Issues - Addressing critical medical imaging challenges including protocol variability (1,251 multi-institutional MRI scans with varying acquisition parameters), image artifacts (noise, motion, intensity inhomogeneities), extreme class imbalance (tumors occupy 5-10% of brain volume requiring weighted loss), privacy regulations (HIPAA/GDPR compliance), clinical consequences (false negatives delay treatment; false positives cause unnecessary interventions), and zero ground-truth leakage requirements.

P5: Extent of Applicable Standards - Compliance with medical imaging formats (NIfTI for research, DICOM for clinical deployment), data privacy regulations (HIPAA, GDPR), BraTS challenge evaluation protocols, reproducibility standards (seed management, deterministic CUDA operations, dependency versioning), and novel integrity auditing procedures where established benchmarks for GNN medical segmentation did not exist.

P6: Level of Uncertainty - High unpredictability from tumor heterogeneity (extreme variation in glioma appearance, shape, size, intensity, texture), scanner variability (different magnetic field strengths, acquisition protocols, vendor hardware), patient factors (motion during scanning, anatomical variations, pathological diversity), architectural uncertainty (optimal graph construction strategy unclear), and generalization risk (training on 1,251 cases may not cover full clinical diversity).

P7: Component Interdependence - Tightly coupled system requiring orchestration: preprocessing impacts feature quality, graph construction defines model input structure, architecture depends on graph properties, training protocol affects convergence, ensemble strategy requires sufficient model diversity, and clinical integration necessitates standardized outputs for PACS systems and visualization tools.

1.7.2 Complex Engineering Activities (A1-A5)

The research demonstrates five categories of complex engineering activities:

A1: Diversity of Resources - Integration of diverse computational and domain resources including BraTS 2021 dataset (1,251 multi-parametric MRI volumes with expert annotations), software libraries (PyTorch Geometric for GNN, SimpleITK for medical imaging, scikit-image for superpixels), consumer-grade hardware (NVIDIA RTX 2060 6GB VRAM demonstrating accessibility), neuroimaging literature (glioma biology, MRI physics), and ML validation frameworks (cross-validation, ablation, ensembles).

A2: Level of Interdisciplinary Interaction - Continuous collaboration across fields: understanding clinical requirements through radiologist consultation for accuracy and speed needs, translating medical knowledge into graph features and network design, ensuring data integrity through validation protocols preventing medical AI pitfalls, interpreting predictions in clinical context (sensitivity to false negatives vs. positives), balancing computer science reproducibility practices with medical privacy and validation standards, and communicating efficiency-accuracy trade-offs to diverse stakeholders.

A3: Innovation and Novelty - Novel technical contributions including first optimized superpixel-GNN integration for 3D medical images with 15 multi-modal features, comprehensive GNN efficiency vs. accuracy trade-off analysis for brain segmentation, characterization of batch size sensitivity in medical imaging GNN tasks (batch 32 optimal, 64 degrades), systematic ablation showing 5-layer architectural sweet spot (6 layers provide no benefit), 15 comprehensive integrity checks addressing reproducibility crisis in medical AI, and demonstration of CNN-competitive

accuracy (92.92%) with $6.9\times$ speedup and $156\times$ parameter reduction.

A4: Societal and Environmental Consequences - Significant implications for healthcare access (enables deployment in developing countries, rural clinics, mobile diagnostic units where volumetric CNNs impractical), clinical efficiency ($6.9\times$ faster inference reduces patient waiting, enables real-time intra-operative guidance), energy sustainability ($156\times$ smaller models reduce carbon footprint at scale), clinical applications (improved segmentation assists diagnosis, treatment planning for surgery and radiation, therapeutic monitoring), error consequences (direct clinical impact requiring careful validation and uncertainty quantification), privacy protection (ethical patient data handling), and global impact (democratizing AI-assisted medical diagnosis worldwide).

A5: Knowledge Application and Self-Directed Learning - Autonomous learning beyond standard curriculum including MRI physics and T1/T2/FLAIR contrasts, glioma pathology and tumor growth patterns, NIfTI structures and skull-stripping algorithms, graph theory message passing and attention mechanisms, patient-level splitting and integrity checks for medical AI validation, GPU memory management and mixed-precision training, and bridging computer science, medical imaging, and clinical practice through effective communication.

1.8 Thesis/Project Organization

This thesis is organized into six chapters that provide a comprehensive and systematic presentation of the research methodology, experimental results, and contributions:

Chapter 1: Introduction presents the research context, motivation, and scope. It introduces brain tumor segmentation challenges, outlines the efficiency-accuracy dilemma in current deep learning methods, and explains why graph neural networks offer a promising alternative. The chapter defines specific research objectives, describes the significance of the work for clinical deployment, and details the complex engineering aspects of the project.

Chapter 2: Literature Review surveys existing approaches in medical image

segmentation, focusing on CNN-based methods (U-Net, V-Net, attention mechanisms) and emerging graph-based techniques. It critically analyzes state-of-the-art performance on BraTS benchmarks, reviews efficiency optimization strategies, discusses data validation concerns in medical AI research, and identifies gaps that motivate the current work.

Chapter 3: Methodology provides detailed technical descriptions of the proposed approach. It documents the BraTS 2021 dataset characteristics, presents the complete pipeline from MRI preprocessing to graph construction, explains the GNN architecture design and training procedures, describes the patient-level stratified cross-validation strategy, and outlines integrity validation protocols that ensure zero data leakage.

Chapter 4: Results and Analysis presents comprehensive experimental findings. It reports 5-fold cross-validation performance ($90.39\% \pm 0.69\%$ Dice), ensemble results (92.92% Dice), efficiency benchmarking against U-Net baselines ($6.9\times$ speedup, $156\times$ parameter reduction), systematic ablation studies validating architectural choices, batch size sensitivity analysis, and qualitative visualizations of segmentation outputs.

Chapter 5: Discussion interprets the results in broader context. It compares performance against state-of-the-art methods, analyzes the implications of efficiency gains for clinical deployment, discusses the significance of architectural findings (5 vs. 6 layers, batch size effects), acknowledges limitations (preprocessing overhead, feature engineering requirements), addresses potential clinical translation barriers, and suggests directions for future work.

Chapter 6: Conclusion synthesizes the key findings and contributions. It summarizes how the research addresses the stated objectives, highlights the main contributions to graph-based medical imaging, emphasizes the practical impact of efficiency improvements, and reflects on lessons learned about balancing accuracy with computational constraints in medical AI applications.

This introduction chapter has established the foundation for a comprehensive

investigation into efficient graph-based brain tumor segmentation. We have motivated the research by highlighting the critical gap between high-performing but computationally expensive volumetric CNNs and the clinical need for efficient, deployable diagnostic tools. The specific objectives outlined provide a roadmap for rigorous validation, from graph construction to ensemble methods, with emphasis on reproducibility and data integrity. By framing this work within complex engineering problem-solving contexts, we underscore the interdisciplinary nature of medical AI research that requires bridging computer science, medical imaging, and clinical practice. The subsequent chapters will demonstrate how graph neural networks can achieve the ambitious goal of maintaining competitive accuracy (92.92% Dice) while delivering substantial efficiency improvements ($6.9\times$ speedup, $156\times$ parameter reduction) suitable for real-world clinical deployment.

Chapter 2

2 Background Study and Literature Review

Brain tumor segmentation has come a long way over the past two decades. Advances in medical imaging and deep learning have pushed the field forward significantly. Accurately outlining tumor boundaries from multi-parametric MRI scans matters for diagnosis, treatment planning, and tracking patients over time. This chapter walks through how automated segmentation methods have evolved—from classical image processing to the CNN revolution, and now toward graph-based approaches. Along the way, we’ll see why computational efficiency and clinical deployment remain persistent challenges, which is exactly what motivates the GNN framework proposed in this thesis.

2.1 Evolution of Medical Image Segmentation

2.1.1 Classical Approaches

Early brain tumor segmentation depended on traditional computer vision techniques: intensity thresholding, region growing, atlas-based methods. These required lots of manual tuning and struggled with how differently tumors can look from patient to patient. The BraTS Challenge, which started in 2012 [7], changed things by standardizing evaluation protocols and providing large-scale annotated datasets. That really accelerated innovation in the field.

2.1.2 The CNN Revolution: U-Net and Its Variants

Fully convolutional networks really transformed medical image segmentation. When Ronneberger et al. [8] introduced U-Net in 2015, it was a game-changer. The symmetric encoder-decoder architecture with skip connections preserved spatial information during upsampling, which proved crucial. U-Net became the go-to standard

for medical segmentation, achieving way better performance than classical methods while needing minimal training data—thanks to aggressive data augmentation.

Taking U-Net to 3D medical imaging, Çiçek et al. [2] extended it with 3D convolutions, which enabled direct processing of multi-slice MRI scans. But there’s a catch: 3D U-Net’s computational requirements scale cubically with volume dimensions. That means you need substantial GPU memory—8GB+ VRAM—and that limits deployment on devices with fewer resources.

Then came nnU-Net from Isensee et al. [5], a self-configuring framework that automatically optimizes preprocessing, architecture, and training hyperparameters for whatever medical segmentation task you throw at it. It achieves state-of-the-art results across diverse datasets. However, the automatic configuration needs extensive computational resources for hyperparameter search, and the final models stay large (80M+ parameters).

2.1.3 Transformer-Based Architectures

Recent works have explored vision transformers for medical segmentation. Hatamizadeh et al. [4] proposed Swin-UNETR, combining Swin Transformer encoders with CNN decoders, achieving 93.8% Dice on BraTS 2021. TransBTS [9] similarly integrates transformer self-attention with 3D convolutions. While these hybrid architectures push accuracy boundaries, they require even larger computational budgets (92M+ parameters, 24GB+ VRAM), exacerbating deployment challenges for clinical settings.

2.2 Graph Neural Networks in Medical Imaging

2.2.1 Fundamentals of Graph Representation Learning

Graph neural networks extend deep learning to non-Euclidean data structures, learning representations through iterative message passing between connected nodes [3, 6]. Hamilton et al. [3] introduced GraphSAGE (Graph Sample and Aggregate), which learns node embeddings by sampling and aggregating features from local

neighborhoods. The inductive learning capability of GraphSAGE enables generalization to unseen graph structures, making it suitable for medical imaging where each patient represents a distinct graph.

2.2.2 Superpixel Segmentation

Superpixel algorithms group perceptually similar pixels into compact regions, providing mid-level representations that preserve object boundaries while reducing dimensionality. Achanta et al. [1] developed Simple Linear Iterative Clustering (SLIC), a computationally efficient method that performs k-means clustering in combined color-spatial space. SLIC generates uniform superpixels with low computational cost (linear in pixel count) and controllable compactness, making it ideal for large-scale medical image preprocessing.

2.2.3 Graph-Based Medical Segmentation: Existing Work

Several studies have explored graph representations for medical imaging:

Hybrid CNN-GNN Approaches: Most existing work combines CNNs for feature extraction with GNNs for spatial reasoning. These hybrid methods retain the computational overhead of volumetric convolutions while adding graph processing costs, yielding limited efficiency gains.

Region Adjacency Graphs: Some methods construct graphs where nodes represent anatomical structures (organs, lobes) and edges encode spatial relationships. While interpretable, these approaches require prior segmentation of regions, introducing preprocessing dependencies.

Point Cloud Representations: Voxel-to-point sampling combined with PointNet-style architectures has been investigated, but point clouds discard spatial structure information that graphs can preserve through explicit edge connectivity.

2.3 Research Gap and Motivation

2.3.1 The Efficiency-Accuracy Dilemma

Despite impressive accuracy numbers, modern deep learning methods hit persistent deployment barriers:

Computational Requirements: State-of-the-art models like nnU-Net, Swin-UNETR, and TransBTS need 68M-92M parameters and 8GB-24GB of GPU memory. That limits deployment to high-end research hardware. Rural clinics, mobile diagnostic units, and developing regions just don't have access to that kind of infrastructure.

Inference Latency: Processing a full $240 \times 240 \times 155$ MRI volume through deep 3D CNNs takes 50-100ms per patient. That's a problem for real-time clinical workflows where radiologists expect immediate feedback.

Energy Consumption: Large models burn through power during inference. As medical AI scales to millions of patients globally, this raises environmental concerns. Sustainable healthcare AI needs architectures that balance accuracy with resource efficiency.

2.3.2 Unexplored GNN Potential

Graph neural networks have proven themselves in social networks, molecular modeling, and recommendation systems. But their application to medical image segmentation? Still pretty new:

Pure GNN Approaches Lacking: Most medical imaging GNN research uses hybrid CNN-GNN designs. *Nobody's rigorously benchmarked a pure, lightweight GNN architecture for efficient brain tumor segmentation.*

Supapixel-GNN Integration Underexplored: Combining superpixel pre-processing ($890 \times$ dimensionality reduction) with inductive graph learning (GraphSAGE) is an unexplored avenue. It could potentially achieve competitive accuracy with dramatic efficiency improvements.

Reproducibility Gap: Medical AI research suffers from reproducibility is-

sues—insufficient documentation, unreported hyperparameters, potential data leakage. Establishing rigorous validation standards (patient-level splits, comprehensive integrity checks) is critical for the field to mature.

2.4 Positioning This Work

This thesis addresses the identified research gap through three key contributions:

1. Pure GNN Architecture: We demonstrate that a lightweight GraphSAGE model (439K parameters) without any convolutional components can achieve 92.92% Dice—approaching state-of-the-art transformers while offering $156\times$ parameter reduction and $6.9\times$ inference speedup.

2. Rigorous Validation Framework: Patient-level stratified cross-validation with 15 comprehensive integrity checks ensures zero data leakage and establishes reproducibility standards for the field.

3. Clinical Deployment Focus: By achieving real-time performance (12.7ms per patient) on consumer-grade hardware (RTX 2060 6GB), we demonstrate practical viability for resource-constrained healthcare settings.

This literature review establishes that while CNNs and transformers have pushed accuracy boundaries for brain tumor segmentation, their computational demands create deployment barriers that limit clinical impact. Graph neural networks, through exploitation of tumor sparsity via superpixel representations, offer a promising yet underexplored paradigm for efficient medical image analysis. The following chapters detail our systematic investigation of this approach, demonstrating that competitive accuracy need not require massive models when architectural design aligns with domain-specific structural properties.

Chapter 3

3 Research Methodology

This chapter lays out the methodology we used to develop and validate the graph-based neural network for brain tumor segmentation. The research follows a systematic pipeline—data preprocessing, graph construction, model architecture design, training, and rigorous validation. Everything’s designed to ensure reproducibility, eliminate data leakage, and enable fair comparison with baseline methods.

Our framework converts 3D multi-parametric MRI volumes into 2D graph representations. Superpixels become nodes, and spatial adjacency defines edges. This cuts computational complexity while preserving essential tumor characteristics through carefully engineered node features. The GNN processes these graphs using message-passing layers that aggregate neighborhood information, enabling both local texture analysis and global structural reasoning. All experimental procedures follow deterministic protocols with comprehensive integrity checks to ensure scientific rigor and clinical reliability.

3.1 Proposed Framework

The framework progresses from raw multi-parametric MRI input through preprocessing, graph construction with superpixel-based dimensionality reduction ($890\times$ compression), feature engineering (15 domain-specific features), GraphSAGE training, patient-level cross-validation, ensemble aggregation, efficiency benchmarking, and comprehensive validation. Each phase incorporates integrity checks to prevent data leakage.

The framework consists of eight interconnected phases that transform raw MRI data into validated tumor segmentation predictions.

Phase 1: Data Preprocessing standardizes the BraTS 2021 dataset through skull-stripping, resampling to isotropic 1mm^3 voxel spacing, Z-score normalization

per modality, and patient-level stratified 5-fold cross-validation split generation. This prevents data leakage.

Phase 2: Graph Construction is where things get interesting. We convert preprocessed 3D volumes into 2D graph representations. For each patient, we extract 200 axial slices using adaptive selection to focus on tumor-containing regions. The SLIC algorithm segments each slice into superpixels (compactness parameter = 10). Each superpixel becomes a graph node with 15 engineered features that capture intensity statistics, spatial information, and shape descriptors. Graph edges connect spatially adjacent superpixels based on boundary sharing.

Phase 3: Feature Engineering extracts domain-specific features without ground-truth leakage. The 15 features include intensity statistics from all four MRI modalities (T1, T1CE, T2, FLAIR), spatial coordinates, and geometric properties. Critically, we don’t use any tumor label information during feature extraction—that would be cheating and would lead to data leakage.

Phase 4: Model Training employs a 5-layer GraphSAGE architecture with 256 hidden dimensions. We train using Adam optimizer (learning rate 0.001), binary cross-entropy loss with class weights to handle imbalance, batch size 32, and 100 epochs with early stopping. Training uses deterministic seed (42) and CUDA deterministic mode for reproducibility.

Phase 5: Cross-Validation implements patient-level stratified 5-fold splits—no patient appears in both training and validation sets. Each fold trains independently with identical hyperparameters, which enables robust performance estimation and statistical significance testing.

Phase 6: Ensemble Methods combine predictions from all 5 folds using soft voting (averaging sigmoid outputs) to leverage model diversity and improve robustness. The ensemble consistently shows 2-3% improvement over individual models.

Phase 7: Efficiency Benchmarking compares inference speed, memory consumption, and model size against 3D U-Net baseline under identical hardware conditions. We measure per-patient latency, GPU memory usage, and parameter count.

Phase 8: Integrity Validation conducts 15 comprehensive audits to ensure

research integrity and reproducibility: (1) feature dimension verification (15 features per node), (2) patient leakage detection across splits, (3) label distribution analysis across folds, (4) normalization consistency checks (Z-score per modality), (5) seed reproducibility verification (deterministic outputs), (6) graph node count validation (10K nodes per patient), (7) edge connectivity verification (planar graph structure), (8) cross-fold data isolation audit (no shared patients), (9) ground truth label integrity (binary mask validity), (10) MRI modality completeness check (all 4 modalities present), (11) slice count validation (200 slices per patient), (12) superpixel boundary consistency, (13) training/validation/test split ratio verification (900/100/251 patients, 72%/8%/20%), (14) model checkpoint consistency, and (15) prediction output format verification.

3.2 Data Set Analysis

This research utilizes the Brain Tumor Segmentation (BraTS) 2021 dataset, the most comprehensive and widely adopted benchmark for evaluating automated brain tumor segmentation algorithms.

Dataset Composition: The BraTS 2021 training set contains 1,251 multi-institutional glioma cases collected from 19 institutions worldwide. Each case includes four co-registered MRI modalities: native T1-weighted (T1), T1-weighted with contrast enhancement (T1CE), T2-weighted (T2), and Fluid Attenuated Inversion Recovery (FLAIR). All images have been pre-processed by BraTS organizers to normalize orientations, resample to 1mm^3 isotropic resolution, and skull-strip to focus on brain tissue.

Ground Truth Annotations: Expert neuroradiologists manually segmented three nested tumor sub-regions: (1) enhancing tumor (ET), (2) tumor core ($\text{TC} = \text{ET} + \text{necrotic core}$), and (3) whole tumor ($\text{WT} = \text{TC} + \text{peritumoral edema}$). For this research, we focus on binary segmentation (tumor vs. non-tumor) by combining all tumor regions into a single positive class.

Data Characteristics: Glioma tumors exhibit high heterogeneity in appearance, size (ranging from small focal lesions to large infiltrative masses), location

(various brain regions), and intensity patterns across modalities. Tumors typically occupy only 5-10% of total brain volume, creating significant class imbalance. The multi-institutional nature introduces scanner variability (different field strengths, acquisition protocols, vendors) that tests generalization capability.

Patient-Level Splitting: Critical for valid medical AI research, we implement patient-level stratified 5-fold cross-validation where each patient’s complete 3D volume appears in exactly one fold. This prevents slice-level leakage where slices from the same patient could appear in both training and testing, artificially inflating performance. Stratification ensures balanced tumor distribution across folds.

Table 1: BraTS 2021 Dataset Statistics and Cross-Validation Split

Characteristic	Value
Total Patients	1,251
MRI Modalities	4 (T1, T1CE, T2, FLAIR)
Image Resolution	$240 \times 240 \times 155$
Voxel Spacing	1 mm ³ (isotropic)
Training Patients (per fold)	900 (72%)
Validation Patients (per fold)	100 (8%)
Test Patients (per fold)	251 (20%)
Average Tumor Volume	5-10% of brain
Ground Truth Labels	Binary (tumor/non-tumor)
Cross-Validation Folds	5 (patient-level stratified)
Total Graph Nodes (per patient)	~10,000 superpixels
Node Features	15 (intensity + spatial + shape)

3.2.1 Data Collection & Preprocessing

Data Acquisition: The BraTS 2021 dataset was obtained through official registration with the BraTS challenge organizers. All data is de-identified and approved for research use, complying with institutional review board (IRB) requirements and patient privacy regulations.

Preprocessing Pipeline: Although BraTS provides pre-processed images, we apply additional standardization:

- **Skull-Stripping:** Remove non-brain tissue to focus computation on relevant regions (already performed by BraTS, verified in our pipeline)

- **Intensity Normalization:** Apply Z-score normalization independently for each MRI modality to standardize intensity distributions across patients: $x_{norm} = (x - \mu)/\sigma$ where μ and σ are computed per modality per patient
- **Slice Selection:** Extract 200 axial slices per patient using adaptive sampling that prioritizes tumor-containing regions while maintaining anatomical coverage
- **Quality Control:** Validate data integrity through automated checks for missing files, corrupted images, inconsistent dimensions, and ground truth label validity

3.2.2 Feature Engineering and Selection

The success of graph-based segmentation depends on node features that capture tumor-discriminative information without introducing ground-truth leakage. We engineered 15 features per superpixel node, organized into three categories:

1. Multi-Modal Intensity Statistics (12 features): For each of the four MRI modalities (T1, T1CE, T2, FLAIR), we compute statistical descriptors over all voxels within each superpixel:

- Mean intensity: captures average signal
- Standard deviation: measures local intensity variation
- Minimum/Maximum intensities: capture intensity range
- Median intensity: robust central tendency

Why these? Glioma tumors exhibit characteristic intensity patterns: hypo-intense in T1, hyper-intense in T2/FLAIR, rim enhancement in T1CE. These statistics let the GNN learn discriminative tumor signatures.

2. Spatial Features (2 features): Normalized (x, y) centroid coordinates of each superpixel provide spatial context, enabling the model to learn anatomical priors about where tumors typically occur.

3. Geometric/Shape Features (1 feature): Superpixel area captures region size, helping distinguish focal lesions from larger tumor masses.

Critical Design Decision: What’s notably absent from our features is tumor label information or tumor ratio within superpixels. Initial experiments included a "tumor_ratio" feature (percentage of tumor pixels in each superpixel), achieving 98% accuracy—but that’s severe data leakage. After removing this feature and retraining, we got the validated 90.39% performance. That ensures fair evaluation and generalization capability.

3.3 Algorithm/ Model Analysis

3.3.1 Graph Construction Algorithm

The graph construction process converts each 3D MRI volume into a 2D graph representation suitable for GNN processing. The procedure consists of five systematic steps:

Step 1: Slice Extraction - Select 200 axial slices from the 3D volume (indices 20-220 of the 240-slice stack) to focus on brain tissue while maintaining computational efficiency.

Step 2: Superpixel Generation - Apply SLIC (Simple Linear Iterative Clustering) algorithm to each slice with parameters: n_segments=50-100 (target number per slice), compactness=10 (balance between spatial and intensity clustering). SLIC groups similar pixels into perceptually coherent regions that often align with anatomical structures.

Node Count Design Choice: We configure SLIC to generate approximately **80-100 superpixels per slice** (parameter experimentation determined this optimal range), resulting in roughly **10,000 total nodes per patient** across all 200 slices. This design balances computational efficiency with spatial resolution—denser superpixel sampling (e.g., 800 nodes/slice yielding 160K total nodes) would preserve finer texture details but increase graph processing cost by 16×, while sparser sampling (e.g., 20 nodes/slice yielding 4K total nodes) would oversimplify tumor boundaries.

Preliminary experiments confirmed that 80-100 nodes/slice provides the optimal accuracy-efficiency trade-off for brain tumor segmentation.

Step 3: Node Feature Extraction - For each superpixel, compute the 15 features described previously by aggregating information from the 4 MRI modalities.

Step 4: Edge Construction - Build graph edges by determining spatial adjacency: two superpixel nodes are connected if they share a boundary in the image space. This creates a planar graph with average degree 4-6.

Step 5: Node Label Assignment - Assign binary labels (tumor/non-tumor) to each node based on majority voting: if >50% of voxels within a superpixel are labeled as tumor in ground truth, the node is labeled positive.

Computational Efficiency: Converting a 3D volume ($240 \times 240 \times 155 = 8.9\text{M}$ voxels) into a graph ($200 \text{ slices} \times \sim 50 \text{ superpixels average} = \sim 10\text{K nodes}$) reduces representation size by **890** $\times$, enabling efficient GNN processing while preserving tumor boundary information through explicit edge connectivity.

3.3.2 GraphSAGE Architecture

We employ GraphSAGE (Graph Sample and Aggregate), a message-passing neural network designed for inductive learning on graphs with node features.

Core Mechanism: GraphSAGE updates node representations by aggregating information from neighborhood nodes through learned transformations:

$$h_{\mathcal{N}(v)}^{(k)} = \frac{1}{|\mathcal{N}(v)|} \sum_{u \in \mathcal{N}(v)} h_u^{(k)} \quad (1)$$

$$h_v^{(k+1)} = \sigma \left(W \cdot \left[h_v^{(k)} \oplus h_{\mathcal{N}(v)}^{(k)} \right] \right) \quad (2)$$

where $h_v^{(k)}$ is the representation of node v at layer k , $\mathcal{N}(v)$ is the neighborhood of v , $h_{\mathcal{N}(v)}^{(k)}$ is the mean-aggregated neighborhood representation, $\oplus$ denotes concatenation, W is a learnable weight matrix, and σ is ReLU activation.

Architecture Specification:

- **Input Layer:** 15 node features

- **Hidden Layers:** 5 GraphSAGE layers with 256 hidden dimensions each
- **Activation:** ReLU activation after each layer
- **Normalization:** Batch normalization for training stability
- **Output Layer:** Single neuron with sigmoid activation for binary classification
- **Total Parameters:** 439,041

Each layer performs message passing (aggregate neighbors $\rightarrow$ concatenate with self $\rightarrow$ linear transform $\rightarrow$ ReLU activation). The 5-layer depth provides 5-hop receptive field sufficient for tumor context aggregation, as validated through ablation studies showing no benefit from 6 layers.

Design Rationale: The 5-layer depth allows message propagation up to 5 hops in the graph, enabling each node to aggregate information from a wide spatial context. Ablation studies confirmed that 6 layers provide no additional benefit. The 256 hidden dimensions balance expressiveness with parameter efficiency.

3.3.3 Training Protocol

Loss Function: Binary Cross-Entropy (BCE) loss with class weights to address extreme class imbalance (tumor nodes $\sim 10\%$, non-tumor $\sim 90\%$):

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N [w_1 \cdot y_i \log(\hat{y}_i) + w_0 \cdot (1 - y_i) \log(1 - \hat{y}_i)] \quad (3)$$

where $w_1 = 9.0$ (weight for tumor class) and $w_0 = 1.0$ (weight for non-tumor class).

Optimization: Adam optimizer with learning rate 0.001, $\beta_1 = 0.9$, $\beta_2 = 0.999$, weight decay 10^{-5} .

Batch Configuration: Batch size 32 (optimal discovered through empirical analysis), gradient accumulation disabled. Larger batches (48, 64) showed performance degradation.

Training Duration: 100 epochs with early stopping (patience=15 epochs) based on validation Dice coefficient. Typical convergence occurs around epoch 50-70.

Hardware: NVIDIA RTX 2060 with 6GB VRAM, demonstrating accessibility on consumer-grade GPUs. Training time: ~ 5 hours per fold.

Reproducibility: Deterministic training with manual seed 42 for Python, NumPy, PyTorch, and CUDA. Environment variables set: CUBLAS_WORKSPACE_CONFIG=:4096:8, CUDNN_DETERMINISTIC=True.

3.3.4 Ensemble Strategy

To improve robustness and performance, we implement soft voting ensemble:

$$P_{ensemble}(tumor) = \frac{1}{5} \sum_{k=1}^5 P_{fold.k}(tumor) \quad (4)$$

where $P_{fold.k}(tumor)$ is the sigmoid output probability from fold k model. Final prediction is tumor if $P_{ensemble} > 0.5$.

Rationale: Each fold trains on different 80% patient subsets, learning complementary patterns. Averaging reduces prediction variance and improves generalization. Our experiments show consistent 2.5% Dice improvement ($90.39\% \rightarrow 92.92\%$).

3.3.5 Evaluation Metrics

Primary Metric: Dice Similarity Coefficient (DSC), the standard metric for medical image segmentation:

$$\text{Dice} = \frac{2 \times |A \cap B|}{|A| + |B|} \quad (5)$$

where A is predicted tumor region and B is ground truth. Dice ranges from 0 (no overlap) to 1 (perfect agreement).

Secondary Metrics:

- **Sensitivity (Recall):** $\frac{TP}{TP+FN}$ - ability to detect actual tumors
- **Specificity:** $\frac{TN}{TN+FP}$ - ability to correctly identify non-tumor regions
- **Precision:** $\frac{TP}{TP+FP}$ - accuracy of positive predictions
- **Inference Time:** Milliseconds per patient volume

- **GPU Memory:** Peak VRAM consumption during inference
- **Model Size:** Total trainable parameters

3.4 Implementation Details

The complete pipeline is implemented in Python using PyTorch and PyTorch Geometric frameworks. Medical image processing utilizes SimpleITK and nibabel for NIfTI file handling, while graph construction employs scikit-image’s SLIC algorithm for superpixel generation. Preprocessed data is stored as compressed .npz files ($4 \text{ modalities} \times 200 \text{ slices}$) for efficient loading during training. Graph representations are structured as PyTorch Geometric Data objects containing node features ($N \times 15$), edge indices ($2 \times E$), and labels ($N \times 1$).

Several implementation choices warrant emphasis for reproducibility: (1) **deterministic training** with manual seed 42 across Python, NumPy, PyTorch, and CUDA environments using CUBLAS_WORKSPACE_CONFIG and CUDNN_DETERMINISTIC flags, (2) **batch size optimization** discovering that batch 32 significantly outperforms larger batches (48, 64) in medical imaging GNN tasks, (3) **comprehensive integrity auditing** with 15 automated validation checks detecting data leakage, patient overlap, and normalization errors, (4) **efficient data structures** converting 8.9M voxel volumes into 10K node graphs ($890\times$ compression) while preserving tumor-discriminative features, and (5) **ensemble post-processing** using soft voting to aggregate sigmoid probabilities from 5 folds before thresholding at 0.5. All experiments run on consumer-grade hardware (NVIDIA RTX 2060, 6GB VRAM) demonstrating accessibility. The complete codebase with documentation is available for reproduction, enabling other researchers to validate our findings and build upon this work.

This methodology chapter has presented a comprehensive and rigorous approach to graph-based brain tumor segmentation. The proposed framework systematically transforms 3D MRI volumes into graph representations through superpixel-based node generation and careful feature engineering that avoids ground-truth leakage.

The GraphSAGE architecture with 5 layers and 256 hidden dimensions provides an optimal balance between expressiveness and parameter efficiency, validated through systematic ablation studies. Patient-level stratified cross-validation ensures statistically valid performance estimation, while the ensemble strategy leverages model diversity for robust predictions. Comprehensive integrity validation with 15 auditing checks (covering data isolation, feature validity, graph structure, reproducibility, and output verification) guarantees reproducibility and eliminates data leakage, addressing critical concerns in medical AI research. The modular pipeline design enables efficient experimentation, adaptation to new datasets, and potential clinical translation. The subsequent chapter will present experimental results demonstrating that this methodology achieves competitive accuracy (92.92% Dice) with substantial efficiency improvements ($6.9\times$ speedup, $156\times$ parameter reduction) suitable for real-world clinical deployment.

Chapter 4

4 Experimental Results and Analysis

This chapter presents comprehensive experimental results that demonstrate how well the graph-based neural network approach works for brain tumor segmentation. The evaluation covers multiple dimensions: cross-validation performance to quantify predictive accuracy, ensemble methods showing robustness, efficiency benchmarking revealing computational advantages, ablation studies validating design choices, and qualitative analysis showcasing segmentation quality. All experiments follow the rigorous methodology from Chapter 3, with deterministic training protocols ensuring reproducibility.

The results show that graph neural networks achieve competitive segmentation accuracy (92.92% Dice with ensemble) while providing substantial efficiency improvements over volumetric CNN baselines— $6.9\times$ faster inference and $156\times$ fewer parameters. Systematic ablation studies confirm that the 5-layer GraphSAGE architecture hits the sweet spot between expressiveness and efficiency. We also discovered something interesting: batch size sensitivity in medical imaging GNN tasks—that provides actionable insights for future research. These findings collectively demonstrate that graph-based approaches offer a viable and practically advantageous alternative to conventional volumetric methods for clinical brain tumor segmentation.

4.1 Environment Setup

All experiments ran under consistent hardware and software configurations to ensure reproducibility and fair comparison.

Hardware Configuration:

- **GPU:** NVIDIA GeForce RTX 2060 with 6GB VRAM
- **CPU:** Multi-core processor with 5 parallel workers for data loading

- **RAM:** 32GB system memory
- **Storage:** SSD with 200GB+ available space

Software Environment:

- **Operating System:** Linux (Ubuntu 20.04+)
- **Python:** 3.12
- **Deep Learning Framework:** PyTorch 2.0.0 with CUDA 11.8
- **Graph Neural Network Library:** PyTorch Geometric 2.3.0
- **Medical Imaging:** SimpleITK 2.2.1, nibabel 5.1.0
- **Scientific Computing:** NumPy 1.24.3, scikit-image 0.21.0

Reproducibility Configuration:

- **Random Seed:** 42 (fixed across Python, NumPy, PyTorch, CUDA)
- **Deterministic Mode:** CUDA deterministic operations enabled
- **Environment Variables:** CUBLAS_WORKSPACE_CONFIG=:4096:8
- **Floating Point:** FP32 (single precision) for all computations

Worth noting: this consumer-grade GPU configuration (RTX 2060, 6GB VRAM) demonstrates the accessibility of our approach. Contrast that with typical medical imaging deep learning research needing high-end GPUs with 24GB+ VRAM. Training time per fold averaged about 5 hours, making the complete 5-fold cross-validation feasible within reasonable computational budgets.

4.2 Cross-Validation Performance

The primary evaluation employed patient-level stratified 5-fold cross-validation on the BraTS 2021 dataset (1,251 patients), ensuring statistically robust performance estimation without data leakage.

4.2.1 Individual Fold Results

Table 2 presents Dice coefficients for each fold, demonstrating consistent performance across different train-validation splits.

Table 2: 5-Fold Cross-Validation Results on BraTS 2021 Dataset

Fold	Train Patients	Val Patients	Dice (%)
Fold 0	1,000	251	90.41
Fold 1	1,000	251	89.62
Fold 2	1,000	251	90.38
Fold 3	1,000	251	91.06
Fold 4	1,000	251	90.50
Mean $\pm$ Std	-	-	90.39 $\pm$ 0.69

Bar plot showing individual fold performance with mean (90.39%) and standard deviation ($\pm 0.69\%$) indicators. Consistent performance across folds demonstrates robust generalization without overfitting to specific patient subsets.

Four key metrics (Dice Coefficient, Accuracy, Sensitivity, Specificity) visualized as boxplots across 5 folds. Tight distributions with minimal outliers demonstrate stable model behavior and robust generalization across different patient subsets. Dice IQR $\leq 1\%$ confirms low inter-fold variance.

Six-panel visualization showing: (top row) training Dice progression, training loss convergence, and validation Dice performance; (bottom row) validation loss tracking, train vs. validation comparison across all folds, and final test performance summary. Consistent convergence patterns across folds demonstrate reproducible training behavior and absence of overfitting. All curves extracted from actual training logs ensuring scientific integrity.

Key Observations:

- **High Average Performance:** Mean Dice of 90.39% demonstrates strong tumor segmentation capability, surpassing typical clinical acceptability thresholds (85-88%).
- **Low Variance:** Standard deviation of 0.69% indicates robust performance across different patient subsets, suggesting good generalization rather than overfitting to specific data characteristics.

- **Consistent Across Folds:** All folds achieve 89-91% range with minimal spread (1.44% between best and worst), confirming that performance is not dependent on particular train-validation splits.
- **Patient-Level Integrity:** Zero patient overlap between training and validation sets verified through comprehensive auditing, ensuring statistical validity.

4.2.2 Secondary Metrics

Beyond Dice coefficient, we evaluate sensitivity, specificity, and precision to provide comprehensive performance characterization.

Table 3: Comprehensive Performance Metrics (Average Across 5 Folds)

Metric	Value (%)
Dice Coefficient	90.39 ± 0.69
Sensitivity (Recall)	91.24 ± 0.82
Specificity	99.12 ± 0.15
Precision	89.67 ± 0.91

Clinical Interpretation:

- **High Sensitivity (91.24%):** The model successfully detects 91.24% of actual tumor regions, critical for clinical screening where missing tumors (false negatives) has severe consequences.
- **Excellent Specificity (99.12%):** Very few non-tumor regions are incorrectly classified as tumor, reducing false alarm rates and maintaining radiologist trust.
- **Balanced Precision (89.67%):** When the model predicts tumor, it is correct 89.67% of the time, demonstrating reliable positive predictions.
- **Trade-off Analysis:** The slight imbalance (sensitivity > precision) indicates the model errs toward over-segmentation rather than under-segmentation, clinically preferable as radiologists can easily dismiss false positives but cannot recover false negatives.

4.3 Ensemble Performance

To explore upper performance bounds and demonstrate robustness through model averaging, we implemented soft voting ensemble combining predictions from all 5 cross-validation folds.

4.3.1 Ensemble Strategy and Results

Method: For each test patient, predictions from all 5 fold models are obtained and averaged at the probability level (before thresholding):

$$P_{ensemble}(tumor) = \frac{1}{5} \sum_{k=1}^5 P_{fold.k}(tumor) \quad (6)$$

Final binary prediction applies threshold 0.5 to the ensemble probability.

Results: The ensemble achieves **92.92% Dice coefficient**, representing a **+2.53% absolute improvement** over the single-model average (90.39%).

Table 4: Single Model vs. Ensemble Performance Comparison

Method	Dice (%)	Improvement
Best Single Model (Fold 3)	91.06	-
Average Single Model	90.39 $\pm$ 0.69	-
Ensemble (5 Models)	92.92	+2.53%

Analysis:

- **Complementary Learning:** Each fold trains on different 72% patient subsets (900 out of 1,251 total), learning diverse patterns and making independent errors. Averaging reduces individual model biases.
- **Variance Reduction:** Ensemble predictions are more stable than individual models, particularly beneficial for edge cases and ambiguous tumor boundaries.
- **Clinical Deployment:** While ensemble requires $5\times$ more inference computation, the per-patient overhead remains acceptable (63.5ms vs 12.7ms) for non-real-time applications like offline diagnosis or treatment planning.

- **Performance Ceiling:** The 92.92% result approaches state-of-the-art performance on BraTS binary segmentation, demonstrating that graph-based methods can match volumetric CNNs in accuracy.

4.4 Efficiency Benchmarking

A key contribution of this research is demonstrating substantial computational efficiency gains over conventional volumetric approaches while maintaining competitive accuracy.

4.4.1 Baseline Comparison

We compare against a standard 3D U-Net baseline (68M parameters) implemented with identical evaluation protocols under the same hardware conditions. This baseline represents a moderately-sized volumetric CNN architecture—larger state-of-the-art models like nnU-Net (80M) and Swin-UNETR (92M) would show even greater parameter disparity, but we select this 68M variant for apple-to-apple comparison as it achieves competitive accuracy (91.2% Dice) representative of standard volumetric segmentation approaches while avoiding confounds from architectural innovations (e.g., transformers, automatic hyperparameter search).

Table 5: Efficiency Comparison: GraphSAGE (Ours) vs. 3D U-Net Baseline

Metric	Our GNN	3D U-Net	Speedup/Reduction
Inference Time (per patient)	12.7 ms	87.8 ms	6.9× faster
Model Parameters	439,041	68M	156× smaller
GPU Memory (inference)	2.1 GB	8.4 GB	4.0× less
Model Size (disk)	1.7 MB	264 MB	155× smaller
Dice Coefficient	90.39%	91.2%	-0.81% (comparable)

Key Findings:

- **6.9× Faster Inference:** GNN processes each patient in 12.7ms vs. U-Net’s 87.8ms, enabling real-time clinical applications (target <50ms easily achieved).

- **156× Parameter Reduction:** 439K parameters vs. 68M enables deployment on resource-constrained devices including mobile platforms, edge computing units, and low-power medical devices.
- **4× Memory Savings:** 2.1GB GPU memory vs. 8.4GB allows running on consumer GPUs (6GB VRAM sufficient) rather than requiring expensive research-grade hardware (24GB+).
- **Minimal Accuracy Trade-off:** Only 0.81% Dice degradation (90.39% vs. 91.2%) represents an excellent efficiency-accuracy balance. With ensemble (92.92%), GNN actually surpasses U-Net baseline.
- **Storage Efficiency:** 1.7MB model size enables easy distribution, version control, and edge deployment compared to 264MB volumetric models.

4.4.2 Clinical Deployment Implications

These efficiency gains have direct practical consequences for clinical translation:

1. **Real-Time Screening:** 12.7ms inference enables interactive diagnostic workflows where radiologists receive immediate segmentation feedback during MRI review sessions.
2. **Resource-Constrained Settings:** Hospitals in developing regions or rural clinics with limited GPU infrastructure can deploy the 439K-parameter model on affordable consumer hardware.
3. **Mobile Diagnostics:** Model size (1.7MB) and memory footprint (2.1GB) permit deployment on high-end mobile devices or portable diagnostic units for telemedicine applications.
4. **Batch Processing:** Lower memory consumption allows processing multiple patients simultaneously, increasing throughput for large-scale screening programs.

5. **Energy Efficiency:** Reduced computational requirements translate to lower power consumption, important for sustainable healthcare AI and battery-powered mobile applications.

4.5 Ablation Studies

Systematic ablation experiments validate architectural design choices and reveal insights about graph neural networks for medical imaging.

4.5.1 Network Depth Analysis

We compare the baseline 5-layer architecture against a 6-layer variant to determine optimal depth.

Table 6: Ablation Study: Network Depth Comparison

Architecture	Test Dice (%)	Parameters	Insight
5 Layers (Baseline)	84.03	439,041	Efficient sweet spot
6 Layers	84.00	571,153	No benefit from depth

Note on Ablation Dice Values: The ablation study reports lower Dice (84.03%) compared to cross-validation (90.39%) because ablation experiments used Fold 0 only with simplified training configurations to enable fair architectural comparisons. The critical finding is the *relative* comparison: 5 layers and 6 layers perform identically ($\sim 84\%$), confirming that additional depth provides no advantage.

Analysis:

- **Optimal Depth Identified:** 5 layers provide sufficient receptive field (5-hop neighborhood aggregation) to capture tumor context without overfitting.
- **Diminishing Returns:** Adding a 6th layer increases parameters by 30% (571K vs 439K) but yields no accuracy improvement (84.00% vs. 84.03%), demonstrating architectural efficiency.
- **Graph Structure Insights:** Unlike deep CNNs that benefit from 50-100+

layers, graph neural networks reach optimal expressiveness at moderate depth due to message-passing mechanisms aggregating wide spatial context efficiently.

- **Practical Implications:** The 5-layer architecture balances model capacity with parameter efficiency, avoiding unnecessary computational overhead while maintaining performance.

4.5.2 Batch Size Sensitivity

During training experiments, we discovered significant batch size sensitivity in medical imaging GNN tasks—a novel empirical finding for the field.

Table 7: Batch Size Impact on Performance

Batch Size	Dice (%)	Performance
32	84-90	Optimal (stable)
48	~86	Slight degradation
64 (via accumulation)	~83	Significant drop

Interpretation:

- **Fine-Grained Details:** Brain tumors exhibit subtle intensity variations and irregular boundaries. Smaller batches provide noisier but more diverse gradient signals that help preserve these fine-grained patterns during training.
- **Large Batch Smoothing:** Larger batches (48, 64) over-smooth gradients, potentially averaging out critical tumor-specific features across diverse patient presentations.
- **Batch 32 as Sweet Spot:** Balances computational efficiency (reasonable GPU utilization) with gradient signal quality (sufficient diversity without excessive noise).
- **Novel Finding:** While batch size effects are known in general deep learning, this specific sensitivity in medical imaging GNN tasks provides actionable guidance for future research in this domain.

4.6 Qualitative Analysis

Beyond quantitative metrics, qualitative examination of segmentation outputs provides insights into model behavior, failure modes, and clinical applicability. Figures ??, ??, and ?? present representative segmentation examples across varying tumor morphologies.

Left to right: T1CE MRI input, ground truth annotation (green), model prediction (red), overlay comparison. The model accurately captures the irregular tumor boundary and correctly identifies the enhancing tumor core.

Complex morphology with multiple tumor components. Model successfully segments primary mass while showing slight over-segmentation in peripheral edema region, consistent with the precision-sensitivity trade-off (89.67% vs. 91.24%).

Small focal lesion demonstrating model sensitivity to subtle tumor presentations. Accurate detection despite small tumor size validates the graph neural network’s ability to aggregate multi-modal intensity features through message passing.

Qualitative Observations:

- **Boundary Accuracy:** Model predictions closely follow irregular tumor boundaries in T1CE images, suggesting effective superpixel-based graph representation preserves spatial precision.
- **Multi-Modal Integration:** Successful segmentation across varying tumor morphologies (large infiltrative masses, small focal lesions) demonstrates that the 15 engineered features effectively capture discriminative patterns from all four MRI modalities (T1, T1CE, T2, FLAIR).
- **Over-Segmentation Bias:** Consistent with sensitivity > precision (91.24% vs. 89.67%), some examples show slight over-segmentation in peri-tumoral edema regions. This clinically preferable bias ensures tumor detection while radiologists can easily correct false positives.
- **Failure Mode Analysis:** Rare failures occur in regions with severe motion artifacts or extreme intensity inhomogeneities, suggesting potential for prepro-

cessing improvements or adversarial training for robustness.

Representative segmentation showing T1CE MRI input with ground truth (green overlay) and model prediction (red overlay). The model accurately captures irregular tumor boundaries and correctly identifies the enhancing tumor core.

Complex morphology with multiple tumor components. Model successfully segments primary mass while showing slight over-segmentation in peripheral edema region, consistent with precision-sensitivity trade-off (89.67% vs. 91.24%). Complex morphology with multiple tumor components. Model successfully segments primary mass while showing slight over-segmentation in peripheral edema region, consistent with precision-sensitivity trade-off (89.67% vs. 91.24%).

Small focal lesion demonstrating model sensitivity to subtle tumor presentations. Accurate detection despite small tumor size validates the graph neural network’s ability to aggregate multi-modal intensity features through message passing.

4.7 Key Findings and Discussions

4.7.1 Primary Contributions Validated

1. Competitive Accuracy with Substantial Efficiency Gains

The 92.92% ensemble Dice demonstrates that graph-based approaches can match volumetric CNN performance on medical segmentation tasks while providing $6.9\times$ speedup and $156\times$ parameter reduction. This challenges the prevailing assumption that 3D convolutions are necessary for medical image analysis.

2. Architectural Optimality Through Ablation

Systematic comparison of 5-layer vs. 6-layer architectures confirms that moderate depth suffices for graph neural networks in medical imaging. This finding enables efficient model design without trial-and-error depth exploration.

3. Batch Size Sensitivity Discovery

The empirical observation that batch 32 significantly outperforms batch 64 in medical imaging GNN tasks (84-90% vs. 83%) provides practical guidance for researchers and highlights the importance of hyperparameter sensitivity analysis in

medical AI.

4. Rigorous Validation Methodology

Patient-level stratified cross-validation with 15 comprehensive integrity checks (feature dimensions, patient overlap, label validity, normalization, seed reproducibility, graph structure, data isolation, modality completeness, slice counts, superpixel consistency, split ratios, checkpoint validation, and output format verification) sets a high standard for reproducible medical AI research, addressing widespread concerns about data leakage in the field.

4.7.2 Limitations and Constraints

1. Preprocessing Overhead

Graph construction (superpixel generation, feature extraction) adds approximately 2-3 seconds per patient preprocessing time. While one-time cost, it impacts real-time clinical workflows requiring immediate results from raw MRI scans.

2. Feature Engineering Dependency

The 15 handcrafted node features require domain knowledge and manual design. Unlike end-to-end learned representations in CNNs, graph-based approaches currently depend on careful feature engineering, though this also enables interpretability.

3. Single-Task Focus

Current implementation performs binary (tumor vs. non-tumor) segmentation. Multi-class segmentation (enhancing tumor, tumor core, edema) would require architectural modifications or multi-task learning frameworks.

4. Performance Gap with State-of-the-Art

While 92.92% is competitive, recent transformer-based methods achieve 93-94% Dice on BraTS. The 1-2% gap represents room for improvement through architectural innovations (e.g., hierarchical graphs, attention mechanisms).

4.8 Comparison with State-of-the-Art Methods

Table 8 compares our approach against recent methods reported on BraTS 2021 binary segmentation.

Table 8: Comparison with State-of-the-Art Methods on BraTS 2021

Method	Dice (%)	Parameters	Key Characteristic
3D U-Net (baseline)	91.2	68M	Standard volumetric CNN
nnU-Net (2021)	92.5	80M	Automatic configuration
Swin-UNETR (2022)	93.8	92M	Vision transformer
TransBTS (2022)	93.2	71M	Hybrid CNN-Transformer
GraphSAGE (Ours - Single)	90.39	439K	Graph-based (156$\times$ smaller)
GraphSAGE (Ours - Ensemble)	92.92	2.2M	5-fold ensemble (31$\times$ smaller)

Comparative Analysis:

- **Accuracy Trade-off:** Our single-model Dice (90.39%) trails state-of-the-art by 2-3%, while ensemble (92.92%) nearly matches best methods. The small accuracy gap is acceptable given efficiency advantages.
- **Parameter Efficiency:** Our 439K parameters are 156 $\times$ smaller than typical CNNs (68-92M), enabling deployment scenarios impossible for volumetric models. Even 5-fold ensemble (2.2M total) remains 31 $\times$ more compact.
- **Inference Speed:** 12.7ms per patient (6.9 $\times$ faster than U-Net baseline) enables real-time applications where transformers (typically slower than CNNs) face deployment challenges.
- **Different Paradigm:** While CNNs/transformers operate on dense 3D grids, our graph-based approach exploits tumor sparsity (5-10% of brain volume), fundamentally different computational strategy.
- **Complementary Strengths:** State-of-the-art transformers excel in accuracy through massive compute; our GNN excels in efficiency through structural

representation. Optimal choice depends on deployment constraints (accuracy-critical vs. resource-constrained).

This chapter has presented comprehensive experimental validation of the proposed graph-based segmentation approach. Cross-validation results ($90.39\% \pm 0.69\%$ Dice) demonstrate robust performance with low variance across patient subsets. Ensemble methods improve accuracy to 92.92% , approaching state-of-the-art while maintaining efficiency advantages. Systematic efficiency benchmarking quantifies $6.9\times$ inference speedup, $156\times$ parameter reduction, and $4\times$ memory savings compared to 3D U-Net baseline—concrete evidence that graph neural networks offer practical deployment advantages. Ablation studies validate the 5-layer architecture as optimal (6 layers provide no benefit) and reveal batch size sensitivity as a novel insight for medical imaging GNN research. Comparison with state-of-the-art methods contextualizes our contributions: while trailing transformers by 1-2% in accuracy, the proposed approach delivers 31-156 $\times$ parameter efficiency suitable for resource-constrained clinical deployment. These results collectively demonstrate that graph-based approaches represent a viable and practically advantageous paradigm for medical image segmentation, challenging conventional assumptions that volumetric processing is necessary for competitive performance.

Chapter 5

5 Standards, Constraints, and Milestones

This chapter addresses the practical, ethical, and operational sides of deploying graph neural networks for clinical brain tumor segmentation. Chapter 4 validated technical performance (92.92% Dice, $6.9\times$ speedup, $156\times$ parameter reduction), but successful clinical translation requires more. We need to think about regulatory standards, computational constraints, ethical implications, and societal impacts. Medical AI systems operate under stringent requirements for safety, privacy, interpretability, and equity—these fundamentally shape how we design and deploy systems.

We'll examine sustainability standards governing medical device software and data management practices, ensuring long-term viability and environmental responsibility. Societal impacts get analyzed through multiple lenses: healthcare accessibility in resource-constrained settings, clinical workflow integration, and potential for diagnostic equity improvements. Ethical considerations address algorithmic fairness, patient privacy, informed consent, and accountability frameworks necessary for responsible AI deployment in high-stakes medical decision-making.

The chapter also discusses technical and operational challenges encountered during development—data heterogeneity, computational limitations, validation complexity. Design constraints from hardware availability, regulatory requirements, and clinical usability needs are explicitly documented. A project timeline with Gantt chart shows the systematic research methodology across data acquisition, model development, validation, and documentation phases.

Understanding these non-technical dimensions is essential for translating research prototypes into clinically viable systems. This chapter provides transparency about real-world constraints and establishes a framework for responsible deployment of AI-assisted brain tumor segmentation in diverse healthcare settings.

5.1 Sustainability Standards

Medical AI systems must adhere to sustainability standards ensuring long-term operational viability, environmental responsibility, and compliance with evolving regulatory frameworks.

5.1.1 Software Development Standards

ISO/IEC 25010 Software Quality Model: Our implementation prioritizes maintainability through modular code architecture, comprehensive documentation, and version-controlled development. The graph construction, training, and inference modules are designed with clear interfaces enabling independent updates without system-wide modifications.

Reproducibility Standards: Following recent calls for reproducible medical AI research (Nature Medicine, 2023), we implement deterministic training protocols with fixed random seeds (seed 42), CUDA deterministic flags, and explicit dependency versioning (PyTorch 2.0.0, PyTorch Geometric 2.3.0). All 15 integrity checks are automated and version-controlled, enabling independent validation by external researchers.

Open Source Licensing: Code released under permissive MIT license encourages collaborative improvement and enables adaptation for diverse clinical contexts, promoting long-term sustainability through community engagement.

5.1.2 Data Management Standards

NIfTI Format Compliance: Strict adherence to Neuroimaging Informatics Technology Initiative (NIfTI) standards ensures interoperability with existing radiology PACS systems and research databases. All preprocessing preserves NIfTI metadata (orientation, voxel spacing, acquisition parameters) for downstream clinical integration.

DICOM Compatibility: While research implementation uses NIfTI format, clinical deployment pathways require DICOM (Digital Imaging and Communica-

tions in Medicine) integration. Graph construction pipeline is designed to accept DICOM inputs through standard conversion libraries (dcm2niix, SimpleITK), ensuring compatibility with hospital imaging infrastructure.

HIPAA and GDPR Compliance: All data handling procedures comply with Health Insurance Portability and Accountability Act (HIPAA) privacy regulations and EU General Data Protection Regulation (GDPR). Patient identifiers are removed during preprocessing; BraTS dataset used for research is already de-identified per institutional review board approval.

5.1.3 Environmental Sustainability

Energy Efficiency: The $156\times$ parameter reduction and $6.9\times$ inference speedup compared to 3D U-Net baseline translate directly to lower energy consumption. Training on consumer-grade RTX 2060 GPU (6GB VRAM) rather than high-end GPUs (24GB+ VRAM) reduces carbon footprint. That’s particularly important as medical AI scales to millions of patients globally.

Compute Optimization: Our graph-based representation exploits tumor sparsity (5-10% of brain volume), processing only relevant anatomical regions rather than dense 3D grids. This architectural choice reduces unnecessary computation, aligning with green AI principles.

Model Longevity: The lightweight architecture (1.7MB model size) extends deployment lifespan on existing hospital hardware. That reduces e-waste from forced hardware upgrades—common with compute-intensive volumetric models.

5.2 Impacts on Society

5.2.1 Healthcare Accessibility

Resource-Constrained Settings: The 439K-parameter model enables brain tumor segmentation in hospitals lacking expensive GPU infrastructure. Rural clinics, community hospitals in developing nations, and mobile diagnostic units can deploy the system on affordable consumer-grade hardware (RTX 2060, \$300-400), democ-

ratizing access to AI-assisted diagnostics previously limited to wealthy research institutions.

Telemedicine Applications: Small model size (1.7MB) and low memory footprint (2.1GB GPU RAM) permit deployment on edge computing devices and potentially high-end mobile platforms. Radiologists conducting remote consultations can provide immediate segmentation analysis without requiring cloud infrastructure or high-bandwidth internet connections.

Batch Screening Programs: Lower computational requirements allow processing multiple patients simultaneously, increasing throughput for population-level screening initiatives in regions with high brain tumor incidence.

5.2.2 Clinical Workflow Integration

Real-Time Assistance: 12.7ms inference time enables interactive workflows where radiologists receive immediate segmentation overlays during MRI review sessions. This real-time feedback can accelerate diagnostic decision-making, reduce cognitive load, and improve clinical efficiency.

Second Opinion Systems: Automated segmentation serves as computational second opinion, highlighting potential tumor regions that may warrant closer expert examination. This is particularly valuable in training settings where junior radiologists benefit from AI guidance.

Treatment Planning Support: Accurate tumor delineation (92.92% Dice) provides quantitative inputs for radiation therapy planning, surgical navigation, and tumor volume monitoring during longitudinal care.

5.2.3 Diagnostic Equity

Reducing Geographic Disparities: By enabling deployment in resource-limited settings, the system can help bridge the diagnostic quality gap between affluent urban hospitals and underserved rural/international communities.

Skill Augmentation: AI assistance may help less-experienced radiologists achieve performance closer to subspecialty experts, particularly important in settings lacking

dedicated neuroradiologists.

Cost Reduction: Computational efficiency translates to lower operational costs (energy, hardware depreciation), potentially making AI-assisted diagnostics economically viable in public health systems with limited budgets.

5.3 Ethics

5.3.1 Algorithmic Fairness and Bias

Dataset Diversity: BraTS 2021 includes multi-institutional data from 19 international sites, partially mitigating single-institution bias. However, representation across racial, ethnic, and socioeconomic groups requires explicit validation before global deployment.

Performance Disparities: Future work must audit model performance across patient subgroups (age, tumor type, scanner manufacturer) to detect potential disparities. Current 5-fold cross-validation provides patient-level generalization but does not stratify by demographic factors.

Bias Mitigation Strategies: Class-balanced training (tumor weight 9.0, non-tumor weight 1.0) addresses extreme imbalance but does not account for subtle population-level variations. Continuous monitoring and retraining on diverse patient cohorts is essential for equitable deployment.

5.3.2 Privacy and Data Protection

De-identification Requirements: All training data must be stripped of Protected Health Information (PHI) per HIPAA standards. Graph-based representation inherently provides some privacy protection—superpixel features are spatially aggregated, making voxel-level reconstruction more difficult than with raw imaging data.

Federated Learning Potential: The lightweight model architecture is well-suited for federated learning paradigms where hospitals collaboratively train models without sharing patient data. Future work could implement secure multi-party

computation for privacy-preserving model updates.

Adversarial Privacy Attacks: Model inversion and membership inference attacks pose theoretical risks. Defensive strategies (differential privacy, model distillation) warrant investigation before clinical deployment.

5.3.3 Clinical Accountability

Decision Support, Not Replacement: The system is designed as a diagnostic assistance tool for radiologists—not an autonomous decision-maker. Final diagnostic and treatment decisions remain under physician responsibility, with AI providing quantitative second opinion.

Interpretability Requirements: While graph structure offers some interpretability (tumor regions correspond to connected superpixel clusters), black-box neural network predictions need explainability methods (attention visualization, saliency maps) for clinical trust and error analysis.

Liability Frameworks: Legal responsibility for AI-assisted diagnoses remains under development globally. Clear documentation of system limitations (1-2% lower accuracy than state-of-the-art transformers, binary segmentation only) is essential for informed clinical use.

5.3.4 Informed Consent

Patient Notification: Patients should be informed when AI systems contribute to diagnostic workflows. Transparency about automated segmentation methods, accuracy rates, and human oversight mechanisms respects patient autonomy.

Opt-Out Mechanisms: Healthcare systems deploying AI assistance should provide options for patients preferring traditional human-only diagnostic processes, acknowledging diverse preferences regarding algorithmic decision involvement.

5.4 Challenges

5.4.1 Technical Challenges

Data Heterogeneity: MRI scans exhibit substantial variability across scanners (GE, Siemens, Philips), acquisition protocols (field strength 1.5T vs. 3T), and image quality. While Z-score normalization addresses intensity variations, geometric distortions and resolution differences remain challenging.

Tumor Diversity: Brain tumors display extreme morphological heterogeneity—from small nodular lesions to large infiltrative masses with irregular boundaries. A single architecture must generalize across this pathological spectrum, requiring careful feature engineering and diverse training data.

Ground Truth Ambiguity: Manual tumor annotations contain inter-rater variability (expert radiologists may disagree on boundaries). Models trained on noisy labels may learn annotation inconsistencies rather than true tumor characteristics.

Graph Construction Overhead: Superpixel segmentation and feature extraction add 2-3 seconds preprocessing per patient. While one-time cost, it impacts real-time clinical workflows requiring immediate results from raw scans.

Batch Size Sensitivity: The discovered sensitivity to batch size (batch 32 optimal, batch 64 degrades to 83%) complicates hyperparameter tuning and transfer to different hardware configurations. This phenomenon requires deeper investigation to establish principled selection guidelines.

5.4.2 Operational Challenges

Clinical Integration Complexity: Integrating AI systems into hospital Picture Archiving and Communication Systems (PACS) requires extensive IT infrastructure work, institutional review board approvals, and clinical workflow redesign—far beyond technical model development.

Regulatory Approval: Medical AI devices require FDA 510(k) clearance (USA), CE marking (Europe), or equivalent regulatory approval before clinical deployment.

This process involves extensive validation studies, clinical trials, and documentation—often taking years and significant financial investment.

Reimbursement Uncertainty: Healthcare reimbursement policies for AI-assisted diagnostics remain unclear in many jurisdictions. Without established billing codes and payment mechanisms, hospital adoption faces financial barriers.

Physician Trust and Adoption: Radiologists may resist AI systems due to concerns about job security, liability, or skepticism about black-box predictions. Building clinical trust requires transparent communication, extensive validation, and demonstration of workflow benefits.

5.4.3 Validation Challenges

Limited Test Data: Despite 1,251 BraTS patients, this represents tiny fraction of global brain tumor population diversity. Generalization to rare tumor subtypes, pediatric cases, or post-treatment imaging requires additional validation cohorts.

Reproducibility Barriers: Many published medical AI results lack code, trained models, or sufficient implementation details for independent replication. Our rigorous documentation addresses this, but field-wide reproducibility crisis remains.

Domain Shift in Deployment: Models validated on research datasets may degrade when deployed in real clinical settings with different scanners, protocols, and patient populations. Continuous performance monitoring and adaptive retraining mechanisms are necessary.

5.5 Constraints

5.5.1 Design Constraints

Clinical Usability Requirements: Segmentation must complete in under 50ms per patient for real-time workflows. Our 12.7ms inference satisfies this constraint with substantial margin, enabling interactive clinical applications.

Accuracy Threshold: Brain tumor segmentation requires minimum 85-90% Dice coefficient for clinical utility, balancing sensitivity (detecting all tumor tissue)

with specificity (avoiding false alarms). Our 92.92% ensemble exceeds this threshold.

Interpretability Needs: Black-box predictions are insufficient for high-stakes medical decisions. Graph structure provides some transparency (tumor corresponds to connected superpixel regions), but further explainability methods (attention visualization) are needed for full clinical acceptance.

Multi-Modality Integration: Brain tumor diagnosis requires synthesizing information from four MRI modalities (T1, T1CE, T2, FLAIR). Graph node features explicitly encode multi-modal intensity statistics, ensuring comprehensive anatomical context.

5.5.2 Component Constraints

Hardware Limitations: Development constrained to consumer-grade RTX 2060 GPU (6GB VRAM), limiting batch sizes and model architectures. This constraint drove architectural efficiency focus but prevented exploration of larger models (transformers requiring 24GB+ VRAM).

Framework Dependencies: PyTorch Geometric library (version 2.3.0) imposes specific data structures (`torch_geometric.data.Data`) and message-passing APIs. While enabling efficient graph operations, framework-specific implementations reduce portability to other platforms (TensorFlow, JAX).

Preprocessing Tools: Dependence on SimpleITK (medical image I/O) and scikit-image (superpixel segmentation) introduces external dependencies. Version compatibility issues and platform-specific bugs required careful dependency management (`requirements.txt` with pinned versions).

Data Format Requirements: BraTS dataset provides NIfTI format files. Clinical deployment requires DICOM compatibility, necessitating format conversion infrastructure and metadata preservation strategies.

5.5.3 Budget Constraints

Zero-Cost Research: Entire project completed without dedicated funding, relying on freely available BraTS dataset, open-source software libraries, and university-

provided GPU access. This constraint necessitated efficient resource utilization and ruled out commercial tools or cloud computing expenses.

Hardware Accessibility: Consumer-grade GPU (RTX 2060, \$300-400 retail) represents accessible price point for academic research. High-end GPUs (RTX 4090, A100) cost \$1,500-10,000+, prohibitive for student projects.

Computational Time Limits: Training 5-fold cross-validation (25 hours total) constrained exploration of larger hyperparameter search spaces. Each experimental iteration required careful planning to maximize insights within limited compute budget.

No Commercial Data Access: Inability to purchase additional annotated datasets limited model validation. Future clinical translation may require commercial datasets or hospital partnerships for diverse patient cohorts.

5.6 Timeline and Gantt Chart

The research project spanned approximately 16 weeks (August 2025 - December 2025), organized into five major phases: dataset acquisition and preprocessing, graph construction pipeline development, model training and validation, comprehensive benchmarking and analysis, and documentation/thesis writing. Table 9 summarizes key milestones and deliverables.

Gantt Chart:

Table 9: Project Timeline and Milestones

Phase	Duration	Key Deliverables
Phase 1: Dataset Acquisition & Preprocessing	Weeks 1-3 (Aug 2025)	BraTS 2021 download, skull-stripping, normalization, 5-fold stratified splits
Phase 2: Graph Construction Pipeline	Weeks 4-7 (Sep 2025)	SLIC superpixel implementation, 15-feature extraction, PyTorch Geometric data conversion, 15 integrity checks
Phase 3: Model Development & Training	Weeks 8-11 (Oct 2025)	GraphSAGE architecture design, 5-fold cross-validation training (5hrs/fold), ensemble implementation
Phase 4: Validation & Benchmarking	Weeks 12-14 (Nov 2025)	Ablation studies (depth, batch size), efficiency benchmarking vs U-Net, qualitative visualization
Phase 5: Documentation & Thesis Writing	Weeks 15-16 (Dec 2025)	Comprehensive documentation, thesis chapters, reproducibility package

Task	Aug	Sep-Oct	Nov	Dec
Dataset Preprocessing	XXXX			
Graph Construction		XXXXXXXX		
Model Training		XXXXXXXX		
Validation Studies			XXXX	
Thesis Writing				XXXX

Key Milestones Achieved:

- **September 15, 2025:** Graph construction pipeline operational with 15 integrity checks passing
- **October 28, 2025:** 5-fold cross-validation training complete, $90.39\% \pm 0.69\%$

Dice achieved

- **November 10, 2025:** Ensemble validation complete, 92.92% Dice confirmed with +2.53% improvement
- **November 25, 2025:** Efficiency benchmarking finalized, $6.9\times$ speedup and $156\times$ parameter reduction quantified
- **December 4, 2025:** All experimental results validated and documented.
- **December 22, 2025:** Paper and thesis submission finalized.

This chapter has examined the practical, ethical, and operational dimensions essential for responsible deployment of graph neural networks in clinical brain tumor segmentation. Sustainability standards ensure long-term system viability through software quality practices, data management compliance (NIfTI, DICOM, HIPAA/GDPR), and environmental responsibility via energy-efficient architectures. Societal impacts span healthcare accessibility in resource-constrained settings, clinical workflow integration for real-time assistance, and potential for diagnostic equity improvements across geographic and skill disparities.

Ethical considerations demand attention to algorithmic fairness (performance auditing across patient subgroups), privacy protection (de-identification, federated learning potential), clinical accountability (decision support with physician oversight), and informed consent frameworks. Technical challenges include data heterogeneity across scanners and protocols, tumor diversity, graph construction overhead, and batch size sensitivity. Operational challenges encompass clinical integration complexity, regulatory approval pathways, reimbursement uncertainty, and building physician trust.

Design constraints shaped system development: real-time performance requirements (12.7ms inference), clinical accuracy thresholds (92.92% exceeds 85-90% minimum), and interpretability needs. Component constraints from hardware limitations (RTX 2060 6GB VRAM), framework dependencies (PyTorch Geometric), and data format requirements influenced architectural choices. Budget constraints (zero

dedicated funding, consumer-grade GPU, limited computational time) necessitated efficient resource utilization and ruled out expensive commercial tools.

The 16-week project timeline demonstrates systematic methodology across dataset preprocessing, graph construction, model training, validation, and documentation phases. Understanding these non-technical dimensions provides essential context for translating research prototypes into clinically viable systems, establishing a framework for responsible AI-assisted medical imaging in diverse healthcare environments.

Chapter 6

6 Conclusion

This thesis shows that graph neural networks provide a viable and efficient way to segment brain tumors, challenging the assumption that you need volumetric 3D convolutions for competitive medical image analysis. The framework we proposed achieves 92.92% Dice coefficient through 5-fold ensemble methods on BraTS 2021 while delivering $6.9\times$ faster inference and $156\times$ fewer parameters compared to 3D U-Net. This enables deployment on consumer-grade hardware and could help democratize access to AI-assisted diagnostics.

Key contributions? First, superpixel-based graph construction that exploits tumor sparsity for $890\times$ dimensionality reduction. Second, validation that a 5-layer GraphSAGE architecture is optimal through ablation studies. Third, discovering batch size sensitivity in medical imaging GNN tasks. And fourth, rigorous validation with patient-level stratified cross-validation and 15 integrity checks addressing data leakage concerns.

Going back to the research question—*Can we achieve competitive accuracy with practical efficiency gains for clinical deployment?*—this thesis provides an affirmative answer. As healthcare AI scales globally, efficiency-aware architectures become essential for equitable healthcare access.

6.1 Limitations

Several limitations are worth acknowledging. First, there’s preprocessing overhead of 2-3 seconds per patient that impacts real-time workflows. Second, handcrafted features require domain expertise and might not generalize to all tumor types. Third, we focused on binary segmentation rather than multi-class delineation. Fourth, validation is limited to BraTS 2021 gliomas—a single dataset. Fifth, there’s still a 1-2% performance gap below transformer methods. And sixth, clinical integration barri-

ers including regulatory approval, hospital IT integration, and real-world validation with radiologists remain unaddressed.

6.2 Future Works and Direction

Several promising directions emerge from this work. Architectural enhancements through attention mechanisms and hybrid CNN-GNN designs could push accuracy higher. Multi-task extensions for multi-class segmentation and survival prediction would increase clinical utility. Preprocessing optimization via differentiable clustering and GPU acceleration could reduce that 2-3 second overhead.

For clinical translation, we'd need radiologist studies, regulatory approval, and multi-institutional trials. Broader applications to other anatomical regions and longitudinal tumor tracking also seem promising.

This thesis establishes graph neural networks as a compelling paradigm for efficient brain tumor segmentation. It demonstrates that competitive medical image analysis doesn't require massive volumetric models. As healthcare AI scales globally, efficiency-aware architectures will be critical for ensuring advanced diagnostics reach diverse populations regardless of resource availability.

References

- [1] Radhakrishna Achanta et al. “SLIC Superpixels Compared to State-of-the-Art Superpixel Methods”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 34.11 (2012), pp. 2274–2282.
- [2] Özgün Çiçek et al. “3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation”. In: *International Conference on Medical Image Computing and Computer-Assisted Intervention*. Springer. 2016, pp. 424–432.
- [3] William L Hamilton, Rex Ying, and Jure Leskovec. “Inductive Representation Learning on Large Graphs”. In: *Advances in Neural Information Processing Systems*. Vol. 30. 2017.
- [4] Ali Hatamizadeh et al. “Swin UNETR: Swin Transformers for Semantic Segmentation of Brain Tumors in MRI Images”. In: *International MICCAI Brainlesion Workshop*. Springer. 2022, pp. 272–284.
- [5] Fabian Isensee et al. “nnU-Net: A Self-Configuring Method for Deep Learning-Based Biomedical Image Segmentation”. In: *Nature Methods* 18.2 (2021), pp. 203–211.
- [6] Thomas N Kipf and Max Welling. “Semi-Supervised Classification with Graph Convolutional Networks”. In: *International Conference on Learning Representations*. 2017.
- [7] Bjoern H Menze, Andras Jakab, Stefan Bauer, et al. “The Multimodal Brain Tumor Image Segmentation Benchmark (BRATS)”. In: *IEEE Transactions on Medical Imaging* 34.10 (2015), pp. 1993–2024.
- [8] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. “U-Net: Convolutional Networks for Biomedical Image Segmentation”. In: *International Conference on Medical Image Computing and Computer-Assisted Intervention*. Springer. 2015, pp. 234–241.
- [9] Wenxuan Wang et al. “TransBTS: Multimodal Brain Tumor Segmentation Using Transformer”. In: *International Conference on Medical Image Computing and Computer-Assisted Intervention*. Springer. 2021, pp. 109–119.